

ArcPy: Solving Large Transportation Analysis Problems

Deelesh Mandloi & Melinda Morang

2020 ESRI DEVELOPER SUMMIT | Palm Springs, CA

Agenda

- Intro/background
- Techniques for solving large problems
- Python script walk-through (solving locally)
- Working with services
- Python script walk-through (solving with a service)

Code and slides:

http://esriurl.com/ds20pysInp

Introduction

ArcGIS Network Analyst Extension for transportation analysis

Coverage

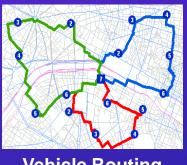


Service Area

Optimization



Location-Allocation



Vehicle Routing Problem

Point-to-point routing



Route



Closest Facility



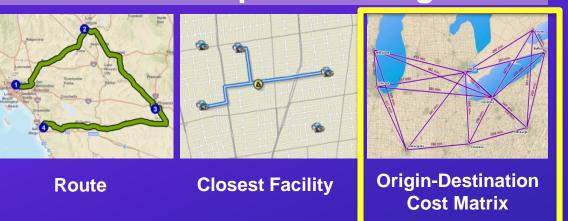
Origin-Destination
Cost Matrix

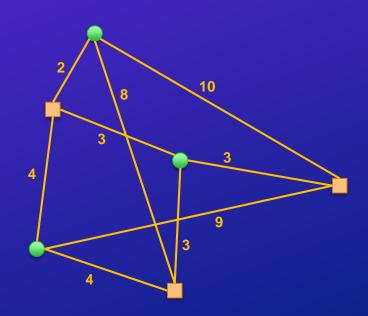
OD Cost Matrix

Calculates the **travel time** or **distance** between a **set of origins** and a **set of destinations**

- Can use a time/distance limit
- Can limit the number of destinations to find

Point-to-point routing





What is a large problem?

- Can't be solved in one calculation
 - Unreasonable calculation time
 - Memory limits
 - Service limits
- Large number of inputs
- Large number of outputs
 - number of origins x number of destinations
 - $-1000 \times 1000 = 1,000,000$

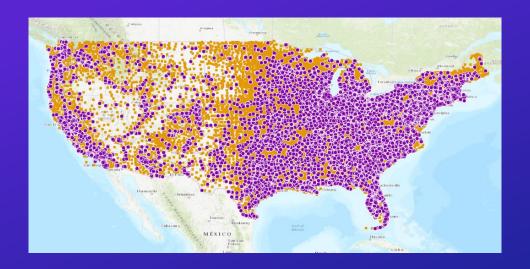
Examples

- Calculate drive time for all patients to all medical clinics within 100 miles
- Calculate the network distance from every parcel to every other parcel

Today's goal

Solve large OD Cost Matrix

- Any number of origins and destinations
- Using local data or a service
- With or without a time/distance limit
- Output a single feature class



What is a local solve and a service solve?

Local solve

- Solving a network analysis on your computer using a network dataset stored on your disk
- Solve uses computing resources of your computer
- Service solve
 - Solving a network analysis using a GIS web service
 - Solve uses computing resources of GIS Server

Which product should I use?



ArcGIS Pro

- 64-bit
- Seamless integration with services
- arcpy.nax module



<u>ArcMap</u>

- 32-bit...runs out of memory easily
- For 64-bit capabilities, must use ArcGIS Server or the 64-bit Background Geoprocessing Extension
- Less simple to use services
- No arcpy.nax module

Network Analysis Workflow with arcpy.nax

- 1. Initialize the analysis object (based on a specific network data source)
- 2. Set the properties for the analysis
- 3. Load the inputs
- 4. Solve the analysis
- 5. Work with the results

arcpy.nax Analysis (Solver) Classes

OriginDestinationCostMatrix

accumulateAttributeNames allowSaveLayerFile defaultDestinationCount defaultImpedanceCutoff distanceUnits ignoreInvalidLocations lineShapeType networkDataSource overrides searchQuery searchTolerance searchToleranceUnits timeOfDay timeUnits timeZone

count()
fieldMappings()
fieldNames()
insertCursor()
load()

solve()

travelMode

Easy-to-use python objects for network analysis

- Added in ArcGIS Pro 2.4
- Analysis class for each solver
 - Set properties
 - Load inputs
 - Solve
- Analysis class for solve results
 - Access outputs

OriginDestinationCostMatrixResult

isPartialSolution solveSucceeded

count()
export()
fieldNames()
saveAsLayerFile()
searchCursor()
solverMessages()

Properties

Methods

Techniques for solving large problems

How to optimize solving a large problem

- Reduce problem size
- Eliminate irrelevant data
- Chunk data
- Spatially sort data
- Solve in parallel
- Pre-calculate network location fields
- Use network dataset layer

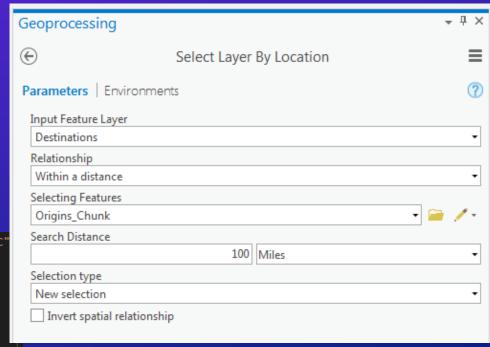
Reduce problem size

- Find only the K nearest
 - defaultDestinationCount property
- Use a time/distance limit
 - Use only destinations within a reasonable straight-line distance of origins
 - defaultImpedanceCutoff property

defaultDestinationCount (Read and Write)	The maximum number of destinations to find per origin. The default is None, which means to find all destinations. The value set in this property can be overridden on a per-origin basis using the TargetDestinationCount field in the input origins.	Integer
defaultImpedanceCutoff (Read and Write)	The impedance value at which to stop searching for destinations from a given origin. If the travel mode used in the analysis uses a time-based impedance attribute, the defaultImpedanceCutoff is interpreted in the units specified in the timeUnits property. If the travel mode used in the analysis uses a distance-based impedance attribute, the defaultImpedanceCutoff is interpreted in the units specified in the distanceUnits property. If the travel mode's impedance attribute is neither time based nor distance based, the defaultImpedanceCutoff value is interpreted in the units of the impedance attribute. The default is None, which means that no cutoff is applied. The defaultImpedanceCutoff can be overridden on a per-origin basis using the Cutoff field in the input origins.	Double

Eliminate irrelevant data

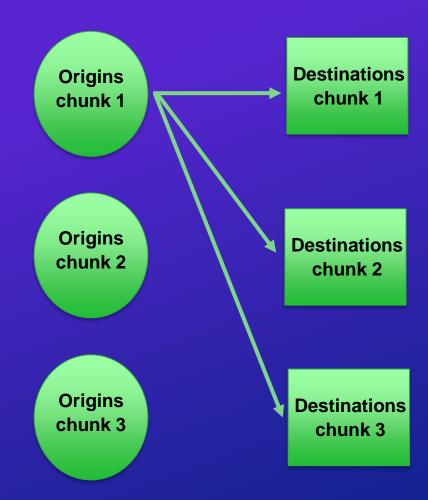
- Remove destinations that aren't within a reasonable straight-line distance of origins prior to running the OD Cost Matrix analysis
- Applies only if you're using a time/distance limit
- Use Select Layer By Location
- Watch out for the case where none are selected



4

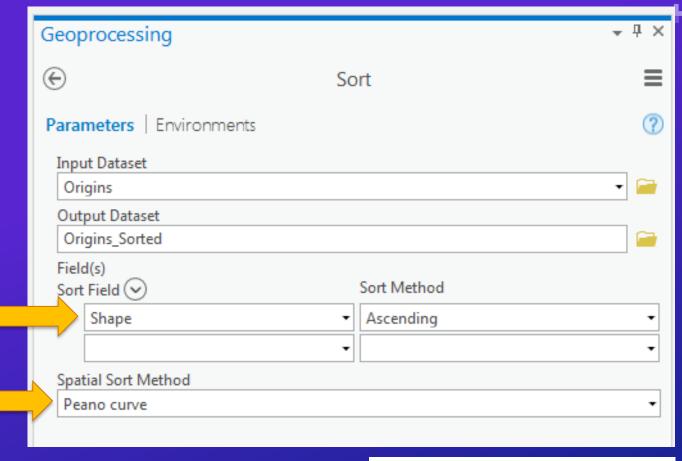
Chunk data

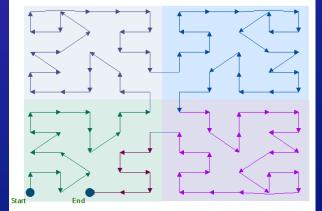
- Break up origins and destinations into chunks of reasonable size
- Iteratively solve each chunk
- Chunk size depends on service limits or memory limits
- Consider number of origins x number of destinations



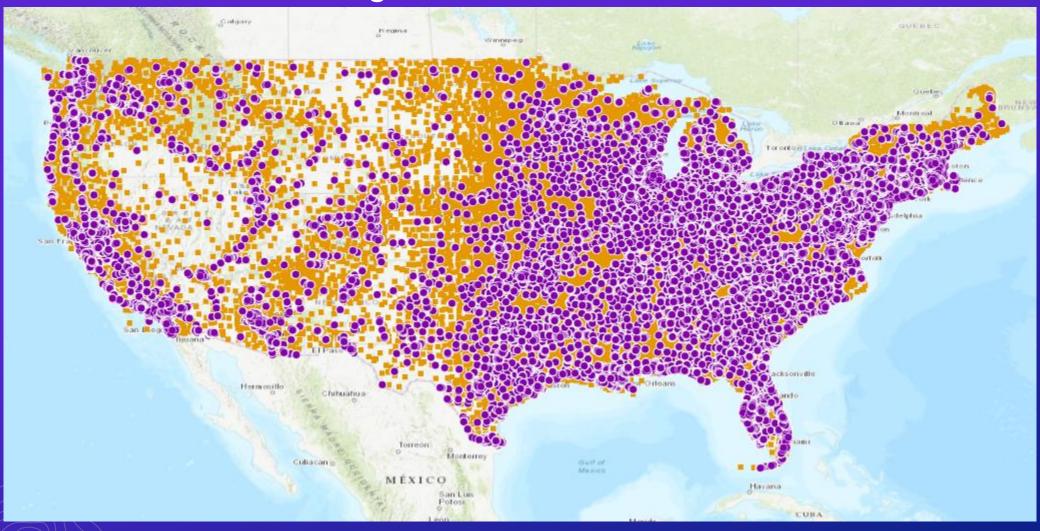
Spatially sort data

- Sort geoprocessing tool
- Sort by Shape field using Peano curve
- Requires Advanced license



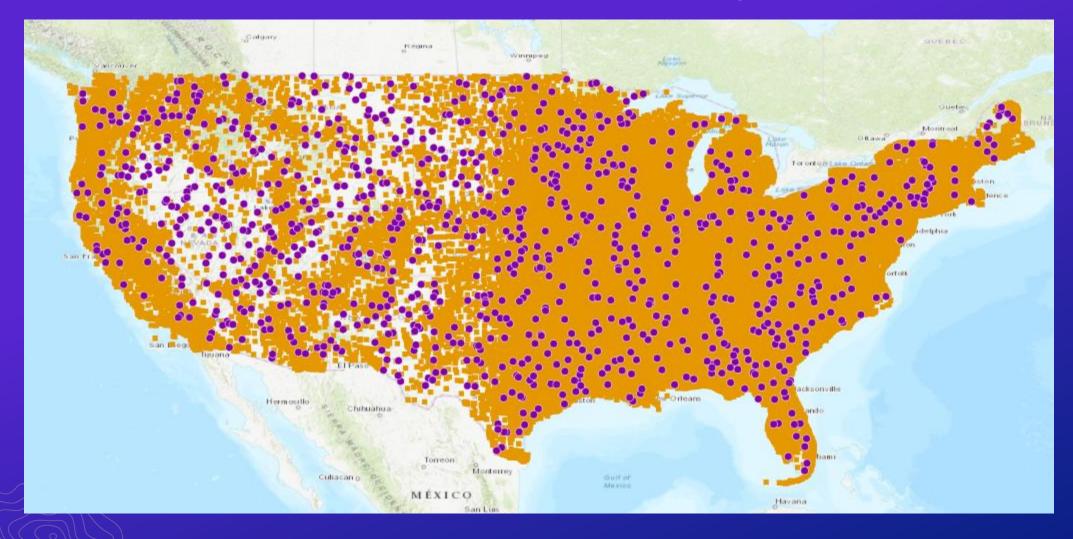


All origins and all destinations

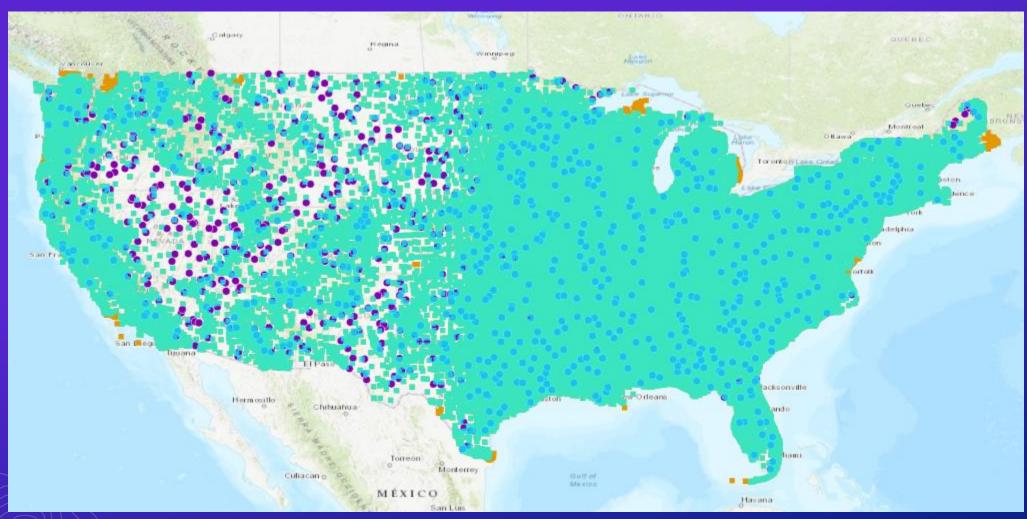


26 million origins , 220k destinations

Chunk of 1000 unsorted origins



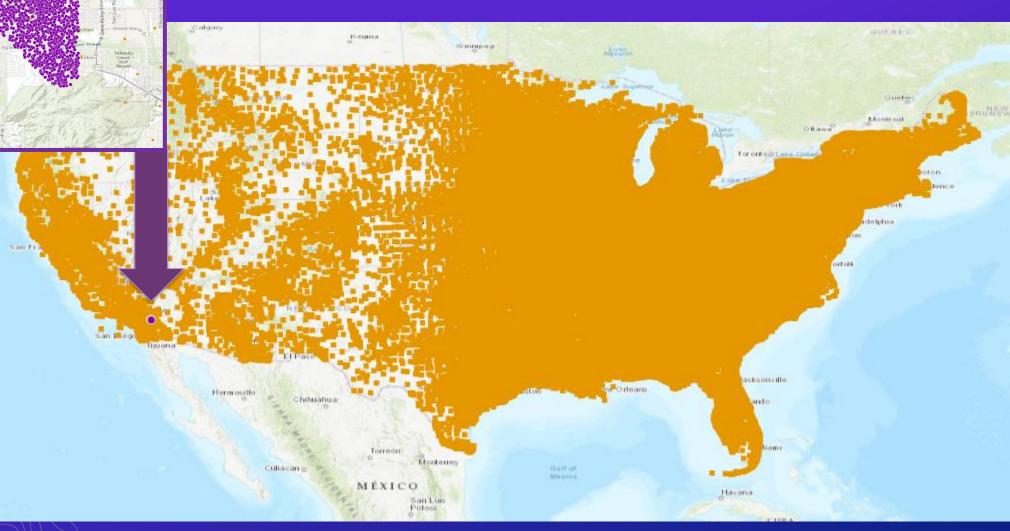
Destinations within 100 miles of 1000 origins



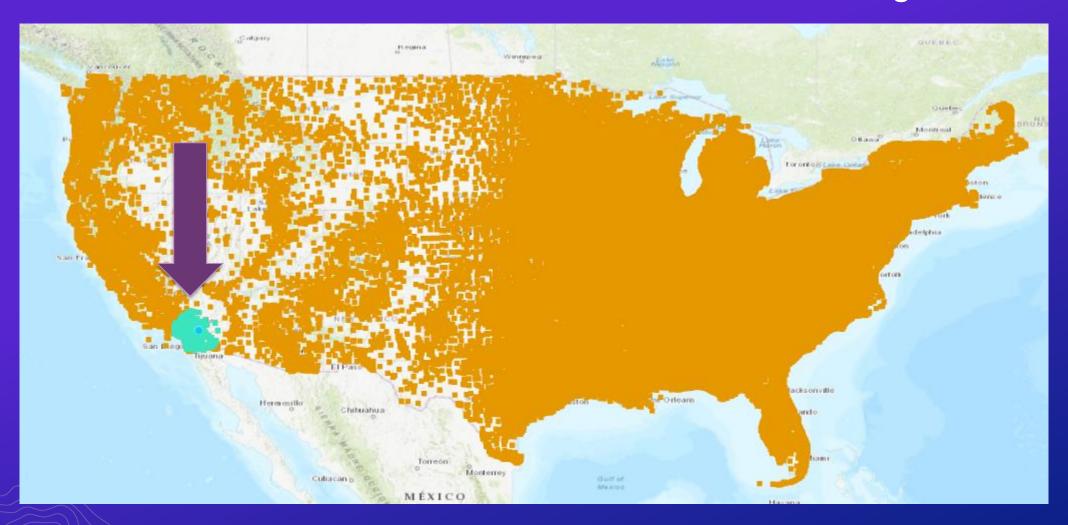
214,449 destinations selected (98%)

E Fances Dr. S. Constant China China

Chunk of 1000 sorted origins



Destinations within 100 miles of 1000 sorted origins



5,653 destinations selected (3%)

Solve in parallel

concurrent.futures

from concurrent import futures

 Multiprocessing: Spin up multiple processes and run solves on multiple cores

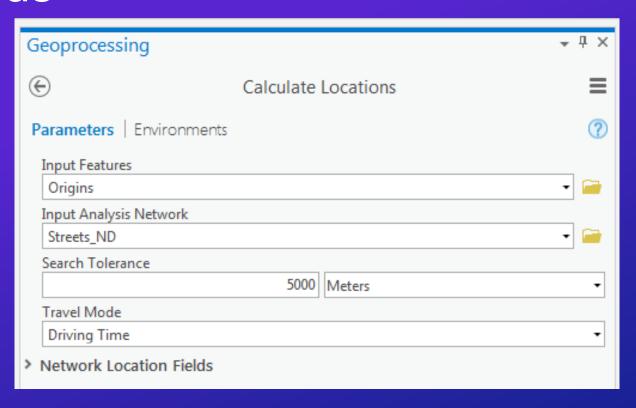
with futures.ProcessPoolExecutor(max workers=inputs["workers"]) as executors:

results = executors.map(solve od cost matrix, inputs iter, ranges)

- Better choice!
- Works well with arcpy
- Easy to run from standalone python
 - If running as a geoprocessing script tool, need to use subprocess module
- Can't write to same gdb from multiple processes
- Multithreading: Use multiple threads in the same process DO NOT USE
 - Not good for CPU-intensive problems in python
 - Does not work with arcpy

Pre-calculate location fields

- Only works for local data (not for services)
- Define how a point snaps to the network
- Calculate them in advance if you're using your points more than once
- Use field mapping in Add Locations to use existing location fields

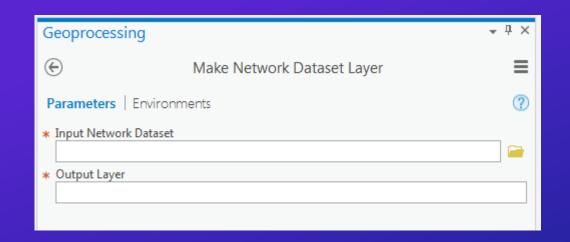




SourceID, SourceOID, PosAlong, SideOfEdge

Use network dataset layer

- Opening from catalog path is slow
- Even slower for licensed data or data on UNC path
- Open once by making a Network
 Dataset Layer; then use the layer
 name (not the layer object)



self.logger.debug("Creating network dataset layer")
arcpy.na.MakeNetworkDatasetLayer(self.network_data_source, self.nds_layer_name)

Outline of today's code

Goal: Solve OD of any size using local data or a service, write it out to a single feature class

- Components:
 - Preprocessing
 - Sorting spatially
 - Calculate network locations
 - Solving
 - Chunking
 - Solving in parallel
 - Post-processing
 - Merging results

Network Analysis Workflow with arcpy.nax

- 1. Initialize the analysis object (based on a specific network data source)
- 2. Set the properties for the analysis
- 3. Load the inputs
- 4. Solve the analysis
- 5. Work with the results

```
# Compute OD cost matrix
od_line_fcs = []
job_folders_to_delete = []
# Run on multiple processes or threads when solving large ODs
if origins count * destinations count > inputs["max od size"]:
    if ODCostMatrix.is_nds_service(inputs["network_data_source"]):
        max workers = os.cpu count() // 2
        pool = futures.ProcessPoolExecutor
        max_workers = (os.cpu_count() // 2) - 1
        pool = futures.ProcessPoolExecutor
    with pool(max_workers=max_workers) as executors:
        results = executors.map(solve od cost matrix, inputs iter, ranges)
        for result in results:
            if result["solveSucceeded"]:
                od_line_fcs.append(result["outputLines"])
                job folders to delete.append(result["jobFolder"])
            else:
                logger.warning("Solve failed for job id %s", result["jobId"])
                logger.debug(result["solveMessages"])
```

Let's look at some code!

(Local solve)

Working with services

Which do I pick?

	Local network dataset	ArcGIS Online service	ArcGIS Enterprise service
Requires your own street data	Yes	No	Yes
Requires Network Analyst extension license	Yes	No	Yes (for Desktop and Server)
Requires service credits	No	Yes	No
Requires ArcGIS Enterprise	No	No	Yes
Number of concurrent processes	Number of CPU cores on your machine	Up to 4	
Analysis limits	Depends on your machine's memory	1000 origins x 1000 destinations	Configurable depending on memory resources of the server

All about credits and service limits!

- Working with ArcGIS Online services requires service credits
 - Credit cost for each service

Origin Destination Cost Matrix	0.0005 credits per input origin and destination pair

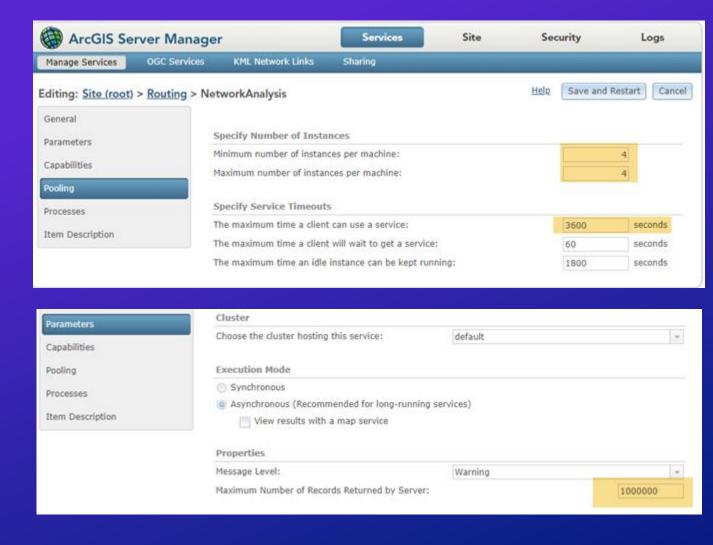
- Each service imposes a limit on the maximum problem size that can be solved in one request. For example, OD service only allows 1,000 origins and 1,000 destinations per request.
 - Service limits for each service
 - OD Cost Matrix service
 - Route service
 - Closest Facility service
 - Service Area service
 - Location-Allocation service
 - Vehicle Routing Problem service

Limit Description	Limit Value
Maximum number of origins	1000
Maximum number of destinations	1000
m number harriers	

Use <u>arcpy.nax.GetWebToolInfo()</u> to retrieve tool limits in a script

How to publish your own routing services to ArcGIS Enterprise

- Use the <u>Publish Routing Services</u> utility
- Special considerations for large problems:
 - Change minimum and maximum service instances to be equal to number of physical CPU cores on your server
 - Set a high value for service usage timeout
 - Set a high value for maximum records returned by server



```
# Compute OD cost matrix
od_line_fcs = []
job_folders_to_delete = []
# Run on multiple processes or threads when solving large ODs
if origins count * destinations count > inputs["max od size"]:
    if ODCostMatrix.is_nds_service(inputs["network_data_source"]):
        max workers = os.cpu count() // 2
        pool = futures.ProcessPoolExecutor
        max_workers = (os.cpu_count() // 2) - 1
        pool = futures.ProcessPoolExecutor
    with pool(max_workers=max_workers) as executors:
        results = executors.map(solve od cost matrix, inputs iter, ranges)
        for result in results:
            if result["solveSucceeded"]:
                od_line_fcs.append(result["outputLines"])
                job folders to delete.append(result["jobFolder"])
            else:
                logger.warning("Solve failed for job id %s", result["jobId"])
                logger.debug(result["solveMessages"])
```

Let's look at some code!

(Service solve)

Wrap-up

- Reduce problem size
- Eliminate irrelevant data
- Chunk data
- Spatially sort data
- Solve in parallel
- Pre-calculate network location fields
- Use network dataset layer

Code and slides:

http://esriurl.com/ds20pyslnp